

```
#importing all necessary libraries
import os
import random
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
plt.style.use("ggplot")
%matplotlib inline

from tqdm import tqdm_notebook, trange
from keras.models import Model, load_model
from keras.layers import Input, BatchNormalization, Activation, Dense, Dropout
from keras.layers.core import Lambda, RepeatVector, Reshape
from keras.layers.convolutional import Conv2D, Conv2DTranspose
from keras.layers.pooling import MaxPooling2D, GlobalMaxPool2D

from keras.optimizers import Adam

from keras.preprocessing.image import ImageDataGenerator, array_to_img, img_to_array, load_img

#set some params , image width, height, border of image
im_width = 128
im_height = 128
border = 5

#print the no. of images in training dataset
ids = next(os.walk("/content/drive/My Drive/pancreas/train/img-panc"))[2]
print("No. of images = ", len(ids))

No. of images = 6029

#initialize array X, y with size of image height, width
X = np.zeros((len(ids), im_height, im_width, 1), dtype=np.float32)
y = np.zeros((len(ids), im_height, im_width, 1), dtype=np.float32)

#tqdm is used to display the progress bar
for n, id_ in tqdm_notebook(enumerate(ids), total=len(ids)):
    try:
        # Load images
        img = load_img("/content/drive/My Drive/pancreas/train/img-panc/{}".format(id_), color_mode="grayscale")
        X_img = img_to_array(img)
        X_img = resize(X_img, (128, 128, 1), mode = 'constant', preserve_range = True)
        # Load masks
        mask = img_to_array(load_img("/content/drive/My Drive/pancreas/train/mask-panc/{}".format(id_), color_mode="grayscale"))

        X[n] = X_img/255.0
        y[n] = mask/255.0
    except:
        continue

/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: TqdmDeprecationWarning: This function will be removed in tqdm==5.0.0
Please use `tqdm.notebook.tqdm` instead of `tqdm.tqdm_notebook`


100% ██████████ 6029/6029 [2:38:33<00:00.158s/it]
```

```
# Split train and valid
X_train, X_valid, y_train, y_valid = train_test_split(X, y, test_size=0.1, random state=42)

# Visualize any random image along with the mask
ix = random.randint(0, len(X_train))
has_mask = y_train[ix].max() > 0 # Pancreas indicator

fig, (ax1, ax2) = plt.subplots(1, 2, figsize = (20, 15))

ax1.imshow(X_train[ix, ..., 0], cmap = 'plasma', interpolation = 'bilinear')
if has_mask: # if Pancreas
    # draw a boundary(contour) in the original image separating Pancreas and non-Pancreas areas
    ax1.contour(y_train[ix].squeeze(), colors = 'k', linewidths = 5, levels = [0.5])
    ax1.set_title('CT-Scan-ABDOMINAL-IMAGING')
ax2.imshow(y_train[ix].squeeze(), cmap = 'gray', interpolation = 'bilinear')
ax2.set_title('MASK-IMG-PANCREAS')
```



```
"""Function to add 2 convolutional layers with the parameters passed to it"""
# first layer
x = Conv2D(filters = n_filters, kernel_size = (kernel_size, kernel_size),\
            kernel_initializer = 'he_normal', padding = 'same')(input_tensor)
if batchnorm:
    x = BatchNormalization()(x)
x = Activation('relu')(x)

# second layer
x = Conv2D(filters = n_filters, kernel_size = (kernel_size, kernel_size),\
            kernel_initializer = 'he_normal', padding = 'same')(input_tensor)
if batchnorm:
    x = BatchNormalization()(x)
x = Activation('relu')(x)

return x

#standard u-net arch
def get_unet_from_img, n_filters = 16, dropout = 0.1, batchnorm = True):
    """Function to define the UNet Model"""
    # Contracting Path
    c1 = conv2d_block(input_img, n_filters * 1, kernel_size = 3, batchnorm = batchnorm)
    p1 = MaxPooling2D((2, 2))(c1)
    p1 = Dropout(dropout)(p1)

    c2 = conv2d_block(p1, n_filters * 2, kernel_size = 3, batchnorm = batchnorm)
    p2 = MaxPooling2D((2, 2))(c2)
    p2 = Dropout(dropout)(p2)

    c3 = conv2d_block(p2, n_filters * 4, kernel_size = 3, batchnorm = batchnorm)
```

```
c4 = conv2d_block(p3, n_filters * 8, kernel_size = 3, batchnorm = batchnorm)
p4 = MaxPooling2D((2, 2))(c4)
p4 = Dropout(dropout)(p4)

c5 = conv2d_block(p4, n_filters = n_filters * 16, kernel_size = 3, batchnorm = batchnorm)

# Expansive Path
u6 = Conv2DTranspose(n_filters * 8, (3, 3), strides = (2, 2), padding = 'same')(c5)
u6 = concatenate([u6, c4])
u6 = Dropout(dropout)(u6)
c6 = conv2d_block(u6, n_filters * 8, kernel_size = 3, batchnorm = batchnorm)

u7 = Conv2DTranspose(n_filters * 4, (3, 3), strides = (2, 2), padding = 'same')(c6)
u7 = concatenate([u7, c3])
u7 = Dropout(dropout)(u7)
c7 = conv2d_block(u7, n_filters * 4, kernel_size = 3, batchnorm = batchnorm)

u8 = Conv2DTranspose(n_filters * 2, (3, 3), strides = (2, 2), padding = 'same')(c7)
u8 = concatenate([u8, c2])
u8 = Dropout(dropout)(u8)
c8 = conv2d_block(u8, n_filters * 2, kernel_size = 3, batchnorm = batchnorm)

u9 = Conv2DTranspose(n_filters * 1, (3, 3), strides = (2, 2), padding = 'same')(c8)
u9 = concatenate([u9, c1])
u9 = Dropout(dropout)(u9)
c9 = conv2d_block(u9, n_filters * 1, kernel_size = 3, batchnorm = batchnorm)

outputs = Conv2D(1, (1, 1), activation='sigmoid')(c9)
model = Model(inputs=[input_img], outputs=[outputs])
```

```
#compiling the model
input_img = Input((im_height, im_width, 1), name='img')
model = get_unet(input_img, n_filters=16, dropout=0.05, batchnorm=True)
model.compile(optimizer=Adam(), loss="binary_crossentropy", metrics=["accuracy"])
```

model.summary()	
conv2d_transpose (Conv2DTranspo	(None, 16, 16, 128) 295040 activation_9[0][0]
concatenate (Concatenate)	(None, 16, 16, 256) 0 conv2d_transpose[0][0]
conv2d_11 (Conv2D)	(None, 16, 16, 128) 295040 activation_7[0][0]
batch_normalization_11 (BatchNo	(None, 16, 16, 128) 512 conv2d_11[0][0]
activation_11 (Activation)	(None, 16, 16, 128) 0 batch_normalization_11[0][0]
conv2d_transpose_1 (Conv2DTrans	(None, 32, 32, 64) 73792 activation_11[0][0]
concatenate_1 (Concatenate)	(None, 32, 32, 128) 0 conv2d_transpose_1[0][0]
conv2d_13 (Conv2D)	(None, 32, 32, 128) 73792 activation_5[0][0]
batch_normalization_13 (BatchNo	(None, 32, 32, 64) 256 conv2d_13[0][0]
activation_13 (Activation)	(None, 32, 32, 64) 0 batch_normalization_13[0][0]
concatenate_2 (Concatenate)	(None, 64, 64, 64) 0 conv2d_transpose_2[0][0]
conv2d_15 (Conv2D)	(None, 64, 64, 32) 18464 activation_3[0][0]
batch_normalization_15 (BatchNo	(None, 64, 64, 32) 128 conv2d_15[0][0]
activation_15 (Activation)	(None, 64, 64, 32) 0 batch_normalization_15[0][0]
conv2d_transpose_3 (Conv2DTrans	(None, 128, 128, 16) 4624 activation_15[0][0]
concatenate_3 (Concatenate)	(None, 128, 128, 32) 0 conv2d_transpose_3[0][0]
conv2d_17 (Conv2D)	(None, 128, 128, 16) 4624 activation_1[0][0]
batch_normalization_17 (BatchNo	(None, 128, 128, 16) 64 conv2d_17[0][0]
activation_17 (Activation)	(None, 128, 128, 16) 0 batch_normalization_17[0][0]
conv2d_18 (Conv2D)	(None, 128, 128, 1) 17 activation_17[0][0]
Total params: 1,179,121	
Trainable params: 1,177,649	
Non-trainable params: 1,472	

```
#creating a checkpoint to save the best model
ReduceLROnPlateau(factor=0.1, patience=5, min_lr=0.00001, verbose=1),
ModelCheckpoint('model-tgs-pancreas.h5', verbose=1, save_best_only=True, save_weights_only=True)
```

```
#fit the data to the model by training it
results = model.fit(X_train, y_train, batch_size=32, epochs=30, callbacks=callbacks,\
                    validation_data=(X_valid, y_valid))

Epoch 00017: val_loss improved from 0.03349 to 0.03197, saving model to model-tgs-pancreas.h5
Epoch 18/30
170/170 [=====] - 63s 4s/step - loss: 0.0235 - accuracy: 0.9750 - val_loss: 0.0326 - val_accuracy: 0.9719

Epoch 00018: val_loss did not improve from 0.03197
Epoch 19/30
170/170 [=====] - 63s 4s/step - loss: 0.0232 - accuracy: 0.9751 - val_loss: 0.0334 - val_accuracy: 0.9721

Epoch 00019: val_loss did not improve from 0.03197
Epoch 20/30
170/170 [=====] - 63s 4s/step - loss: 0.0227 - accuracy: 0.9752 - val_loss: 0.0345 - val_accuracy: 0.9719

Epoch 00020: val_loss did not improve from 0.03197
Epoch 21/30
170/170 [=====] - 63s 4s/step - loss: 0.0232 - accuracy: 0.9748 - val_loss: 0.0346 - val_accuracy: 0.9717

Epoch 00021: val_loss did not improve from 0.03197
Epoch 22/30
170/170 [=====] - 63s 4s/step - loss: 0.0222 - accuracy: 0.9753 - val_loss: 0.0332 - val_accuracy: 0.9720

Epoch 00022: ReduceLROnPlateau reducing learning rate to 0.0001000000474974514e-05.
Epoch 23/30
170/170 [=====] - 63s 4s/step - loss: 0.0223 - accuracy: 0.9754 - val_loss: 0.0332 - val_accuracy: 0.9720

Epoch 00023: val_loss improved from 0.03197 to 0.03185, saving model to model-tgs-pancreas.h5
Epoch 24/30
170/170 [=====] - 63s 4s/step - loss: 0.0196 - accuracy: 0.9760 - val_loss: 0.0319 - val_accuracy: 0.9727

Epoch 00024: val_loss did not improve from 0.03185
Epoch 25/30
170/170 [=====] - 63s 4s/step - loss: 0.0193 - accuracy: 0.9761 - val_loss: 0.0322 - val_accuracy: 0.9727

Epoch 00025: val_loss did not improve from 0.03185
Epoch 26/30
170/170 [=====] - 63s 4s/step - loss: 0.0188 - accuracy: 0.9764 - val_loss: 0.0320 - val_accuracy: 0.9727

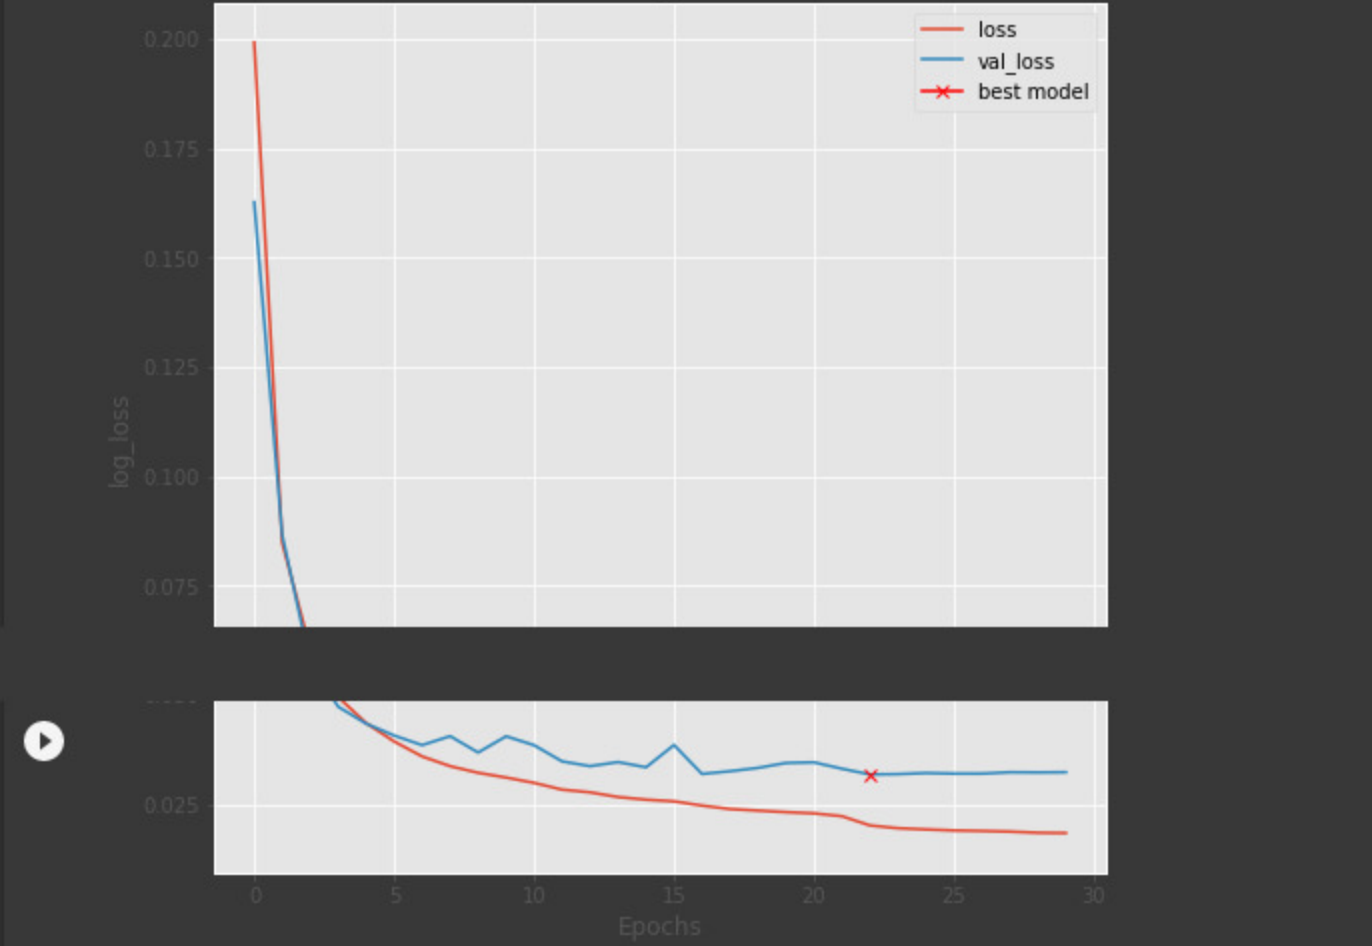
Epoch 00026: val_loss did not improve from 0.03185
Epoch 27/30
170/170 [=====] - 63s 4s/step - loss: 0.0188 - accuracy: 0.9763 - val_loss: 0.0320 - val_accuracy: 0.9727

Epoch 00027: val_loss did not improve from 0.03185
Epoch 28/30
170/170 [=====] - 63s 4s/step - loss: 0.0186 - accuracy: 0.9764 - val_loss: 0.0324 - val_accuracy: 0.9727

Epoch 00028: ReduceLROnPlateau reducing learning rate to 1.00000000474974514e-05.
Epoch 29/30
170/170 [=====] - 63s 4s/step - loss: 0.0187 - accuracy: 0.9761 - val_loss: 0.0323 - val_accuracy: 0.9727

Epoch 00029: val_loss did not improve from 0.03185
Epoch 30/30
170/170 [=====] - 63s 4s/step - loss: 0.0183 - accuracy: 0.9764 - val_loss: 0.0324 - val_accuracy: 0.9727

Epoch 00030: val_loss did not improve from 0.03185
```



```
#saving the model to access frequently
model.save('drive/My Drive/pancreas/train/pancreas_model.h5')

#loading the model to use it in the process
model = load_model('drive/My Drive/pancreas/train/pancreas_model.h5')

# load the best model
model.load_weights('drive/My Drive/pancreas/train/model-tgs-pancreas.h5')

# Evaluate on validation set (this must be equals to the best log_loss)
model.evaluate(X_valid, y_valid, verbose=1)

19/19 [=====] - 16s 815ms/step - loss: 0.0319 - accuracy: 0.9726
[0.0185439109802246, 0.972599089145604]

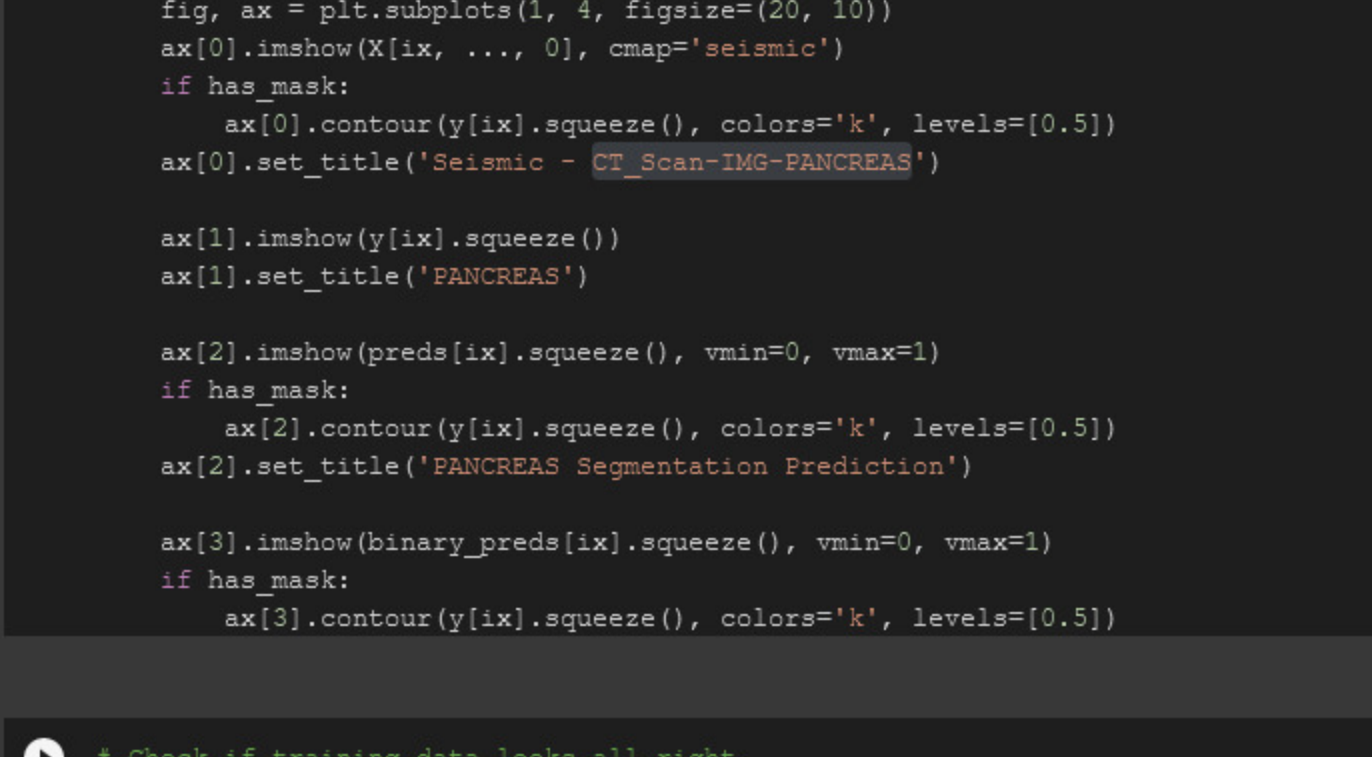
# Predict on train, val and test
preds_train = model.predict(X_train, verbose=1)
preds_val = model.predict(X_valid, verbose=1)

170/170 [=====] - 140s 823ms/step

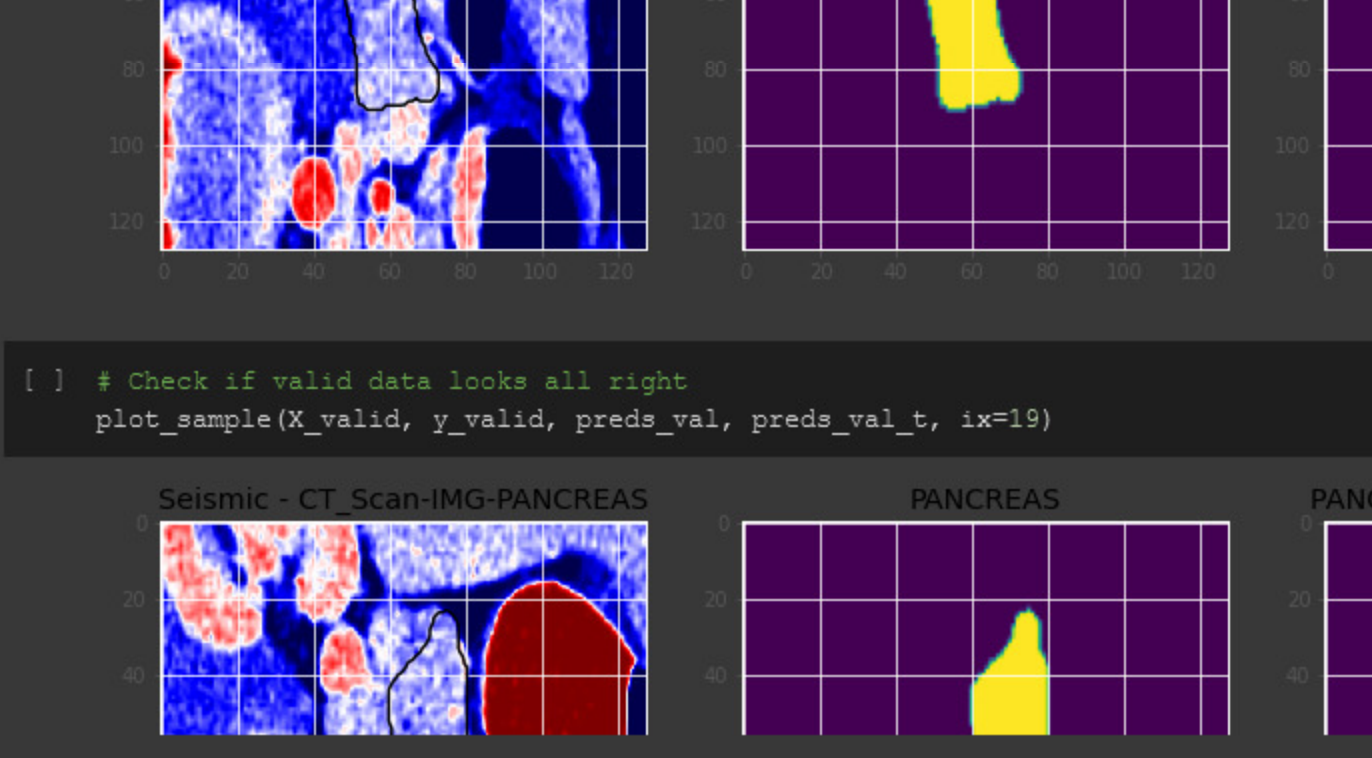
# Threshold predictions
preds_train_t = (preds_train > 0.5).astype(np.uint8)
preds_val_t = (preds_val > 0.5).astype(np.uint8)

#function to plot the data and to check the predictions.
def plot_sample(X, y, preds, binary_preds, ix=None):
    """function to plot the results"""
    if ix is None:
        ix = random.randint(0, len(X))
    has_mask = y[ix].max() > 0

    fig, ax = plt.subplots(1, 4, figsize=(20, 10))
    ax[0].imshow(X[ix, ..., 0], cmap='seismic')
    if has_mask:
        ax[0].contour(y[ix].squeeze(), colors='k', levels=[0.5])
        ax[0].set_title('Seismic - CT_Scan-IMG-PANCREAS')
    ax[1].imshow(y[ix].squeeze())
    ax[1].set_title('PANCREAS')
    ax[2].imshow(preds[ix].squeeze(), vmin=0, vmax=1)
    if has_mask:
        ax[2].contour(y[ix].squeeze(), colors='k', levels=[0.5])
        ax[2].set_title('PANCREAS Segmentation Prediction')
    ax[3].imshow(binary_preds[ix].squeeze(), vmin=0, vmax=1)
    if has_mask:
        ax[3].contour(y[ix].squeeze(), colors='k', levels=[0.5])
```



```
# Check if training data looks all right
plot_sample(X_train, y_train, preds_train, preds_train_t, ix=14)
```



```
#ploting a random sample from validation set
plot_sample(X_valid, y_valid, preds_val, preds_val_t)
```

